

Capstone Project - 2

Team 2 : Loan Default Prediction

Let's Catch The Defaulters

1. Defining problem statement
2. EDA and feature engineering
3. Feature Selection
4. Preparing dataset for modeling
5. Model Validation and Selection





The Dilemma

Lending Club is the world's largest online marketplace connecting borrowers and investors. An inevitable outcome of lending is default by borrowers. The idea of this tutorial is to create a predictive model that identifies applicants who are relatively risky for a loan. In order to accomplish this, I organized the whole series into four parts as follows:


How Lending Club Works



Borrowers apply for loans.
Investors open an account.



Borrowers get funded.
Investors build a portfolio.

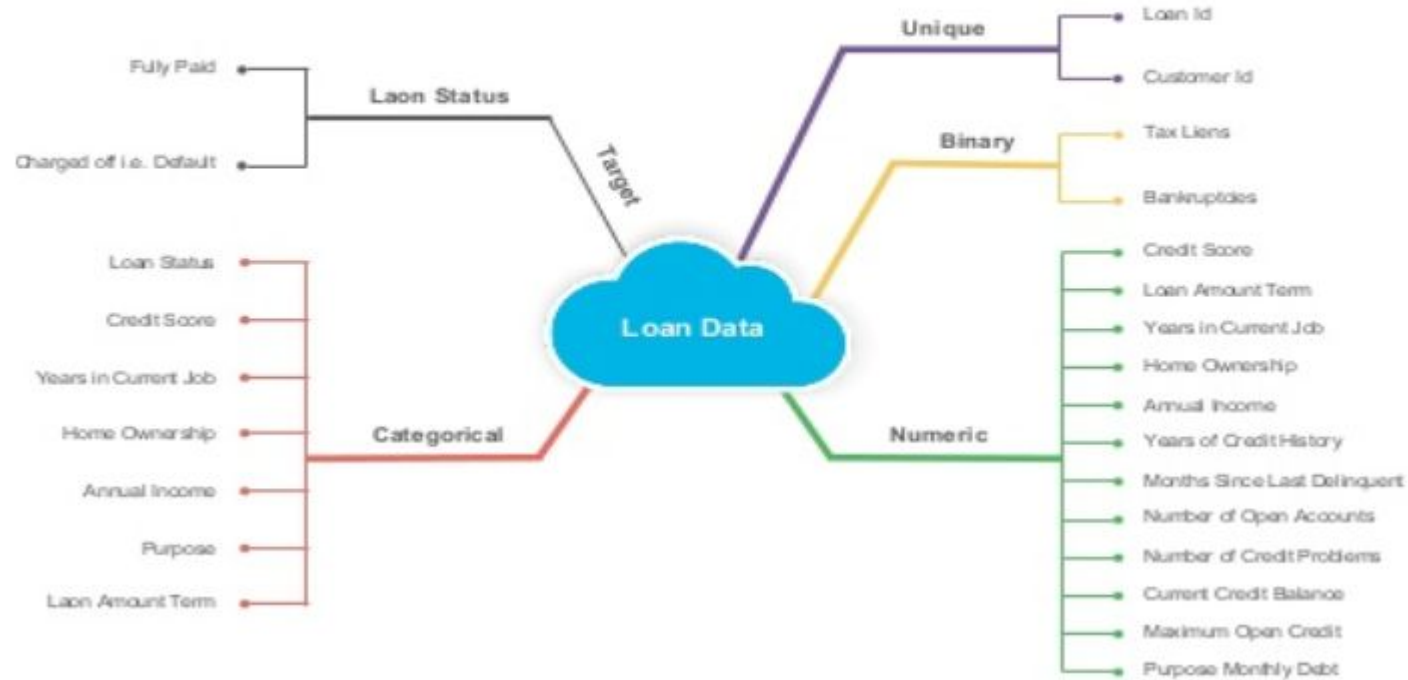


Borrowers repay automatically.
Investors earn & reinvest.

Data Pipeline

- **Data processing-1**: In this first part we've removed unnecessary features. There were nearly Many columns with all Null values.
- **Data processing-2**: In this part, we manually go through each features selected from part 1. This is the most time-consuming part, but worth it for a better model.
- **EDA**: In in this part, we do some exploratory data analysis (EDA) on the features selected in part-1 and 2.
- **Create a model**: Finally, In this last but not the last part, we create models. Creating a model is also not an easy task. It's also an iterative process. we show how to start with a with a simple model, then slowly add complexity for better performance.

Data Summary



Data Summary

loan_amnt: The amount the borrower promises to repay, as set forth in the loan contract. The loan amount may exceed the original amount requested by the borrower if he or she elects to include points and other upfront costs in the loan.

Purpose: The purpose of the loan is used by the lender to make decisions on the risk and may even impact the interest rate that is offered. For example, if an applicant is refinancing a mortgage after having taken some cash out, the lender might consider that an increase in risk and increase the interest rate that is offered or add additional conditions. Loan purpose is important to the process of obtaining mortgages or business loans that are connected with specific types of business activities.

Data Summary

Dti: Your debt-to-income ratio is all your monthly debt payments divided by your gross monthly income. This number is one way lenders measure your ability to manage the monthly payments to repay the money you plan to borrow.


Fico_range: A FICO score is a credit score created by the Fair Isaac Corporation (FICO).¹ Lenders use borrowers' FICO scores along with other details on borrowers' credit reports to assess credit risk and determine whether to extend credit. FICO scores take into account data in five areas to determine creditworthiness: payment history, current level of indebtedness, types of credit used, length of credit history, and new credit accounts.

Data Summary

Delinq_amnt : Delinquent describes something or someone who fails to accomplish that which is required by law, duty, or contractual agreement, such as the failure to make a required payment or perform a particular action.

Int_rate: An interest rate is defined as the proportion of an amount loaned which a lender charges as interest to the borrower, normally expressed as an annual percentage. It is the rate a bank or other lender charges to borrow its money, or the rate a bank pays its savers for keeping money in an account.

Define Dependent Variable



Define Dependent Variable

Fully Paid->Loan has been fully paid off.

Charged off->Loan for which there is no longer a reasonable expectation of further payments.

Does not meet the credit policy. Status:Fully Paid--> While the loan was paid off, the loan application today would no longer meet the credit policy and wouldn't be approved on to the marketplace.

Does not meet the credit policy. Status:charged off->While the loan was charged off, the loan application today would no longer meet the credit policy and wouldn't be approved on to the marketplace.


Current->Loan is up to date on current payments.,

In Grace period->The loan is past due but still in the grace period of 15 days.

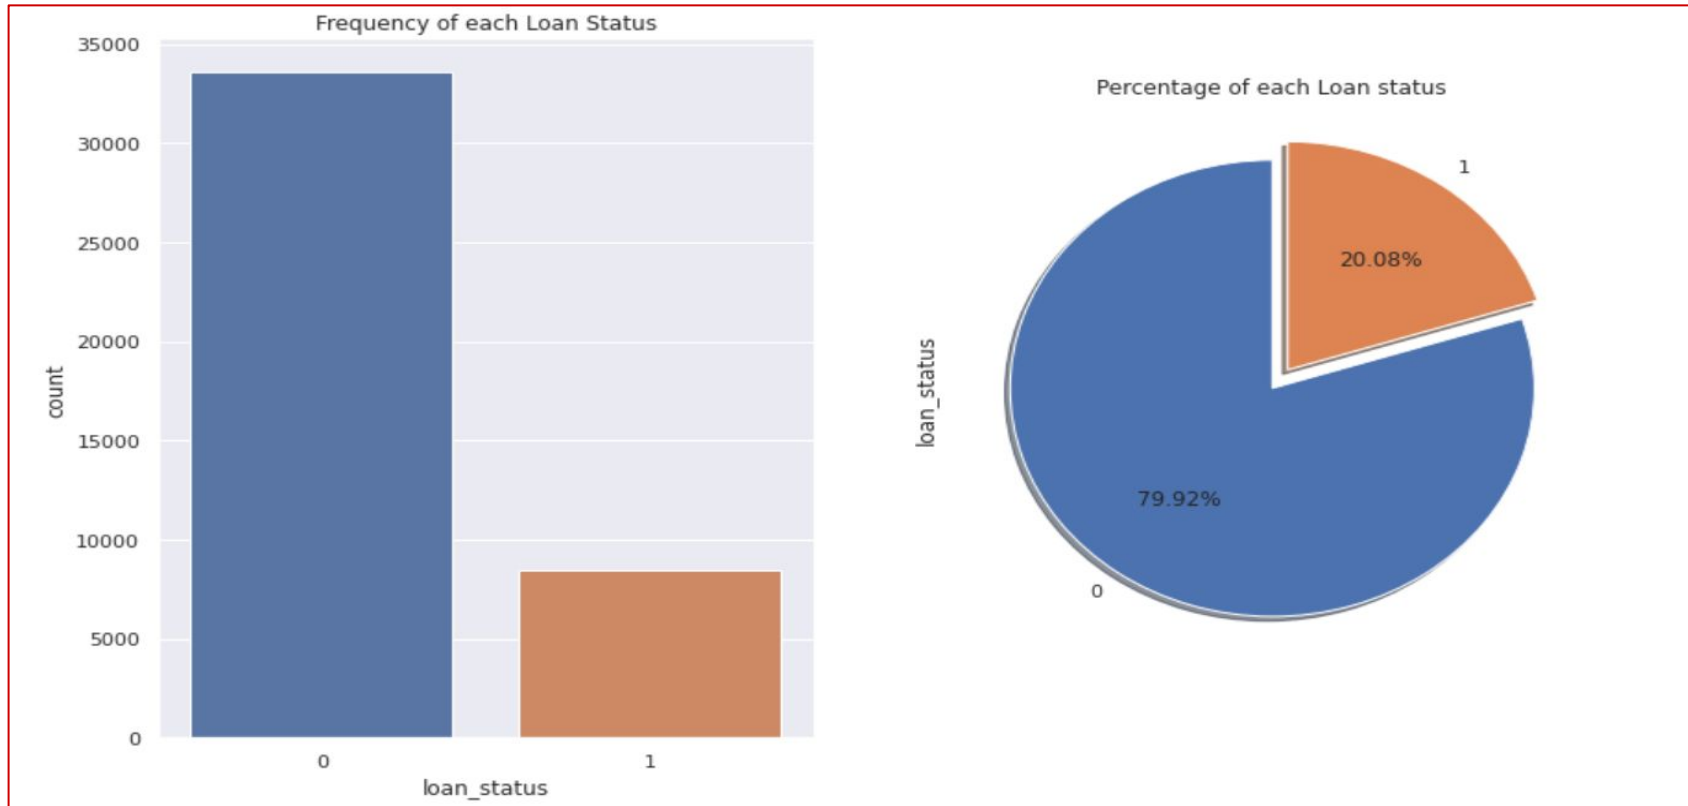
Late(30-120)->Loan hasn't been paid in 31 to 120 days (late on the current payment).

Late(16-30)->Loan hasn't been paid in 16 to 30 days (late on the current payment).

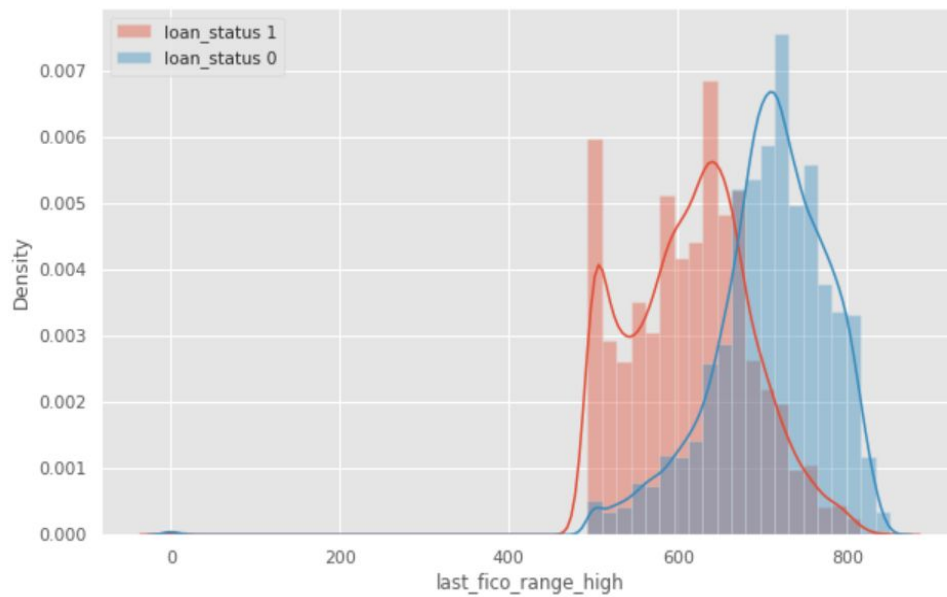

Default->Loan is defaulted on and no payment has been made for more than 121 days.




Define Dependent Variable



EDA




EDA Cont..




EDA cont..


Number of defaulted loans per state



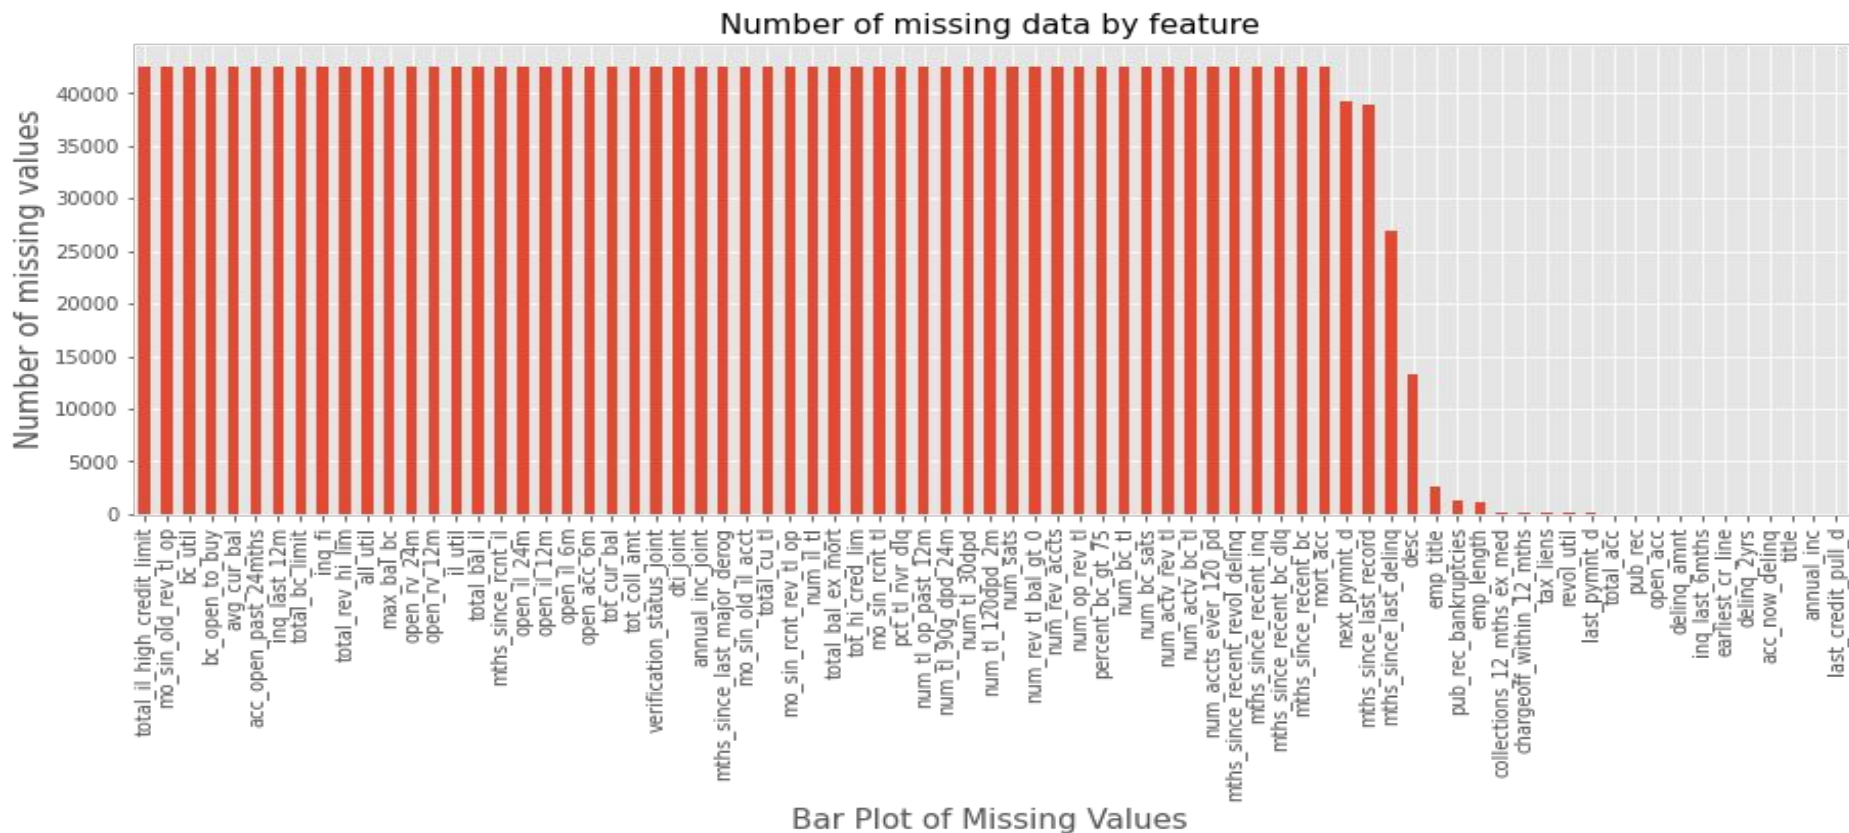
Number of not-defaulted loans per state



EDA cont..




EDA cont..



EDA cont.. (Implement KNN Imputer for missing value)

- Imputation for completing missing values using k-Nearest Neighbors. Each sample's missing values are imputed using the mean value from $n_neighbors$ nearest neighbors found in the training set. Two samples are close if the features that neither is missing are close.
- The KNN Imputer class provides imputation for filling in missing values using the k-Nearest Neighbors approach. By default, a euclidean distance metric that supports missing values, `nan_euclidean_distances`, is used to find the nearest neighbors. Each missing feature is imputed using values from $n_neighbors$ nearest neighbors that have a value for the feature. The feature of the neighbors are averaged uniformly or weighted by distance to each neighbor.



Preparing dataset for modeling

	loan_amnt	term	int_rate	installment	grade	emp_length	annual_inc	purpose	addr_state	dti	delinq_2yr
0	5000.0	0.0	10.65	162.87	1.0	10.0	24000.0	1.0	3.0	27.65	0.0
1	2500.0	1.0	15.27	59.83	2.0	1.0	30000.0	0.0	10.0	1.00	0.0
2	2400.0	0.0	15.96	84.33	2.0	10.0	12252.0	11.0	14.0	8.72	0.0
3	10000.0	0.0	13.49	339.31	2.0	10.0	49200.0	9.0	4.0	20.00	0.0
4	5000.0	0.0	7.90	156.46	0.0	3.0	36000.0	13.0	3.0	11.20	0.0
5	7000.0	1.0	15.96	170.08	2.0	8.0	47004.0	2.0	27.0	23.51	0.0
6	3000.0	0.0	18.64	109.43	4.0	9.0	48000.0	0.0	4.0	5.35	0.0
7	5600.0	1.0	21.28	152.39	5.0	4.0	40000.0	11.0	4.0	5.55	0.0
8	5375.0	1.0	12.69	121.45	1.0	1.0	15000.0	9.0	42.0	18.08	0.0
9	6500.0	1.0	14.65	153.45	2.0	5.0	72000.0	2.0	3.0	16.12	0.0
10	12000.0	0.0	12.69	402.54	1.0	10.0	75000.0	2.0	4.0	10.78	0.0
11	9000.0	0.0	13.49	305.38	2.0	1.0	30000.0	2.0	44.0	10.08	0.0
12	3000.0	0.0	9.91	96.68	1.0	3.0	15000.0	1.0	14.0	12.56	0.0
13	10000.0	0.0	10.65	325.74	1.0	3.0	100000.0	9.0	4.0	7.06	0.0
14	1000.0	0.0	16.29	35.31	3.0	1.0	28000.0	2.0	24.0	20.31	0.0

Task:-

Train Set:-

Test set:-

Response:-

Model Validation & Selection(Baseline Model)

Logistic Regression

test_accuracy	0.832096
recall_test	0.330956
precision_test	0.660969
f1_test	0.441065
auc_test	0.644235
cm_test	[[8046, 357], [1407, 696]]
train_accuracy	0.831165
precision_train	0.662284
recall_train	0.32607
f1_train	0.436991
auc_train	0.642128
cm_train	[[24130, 1053], [4268, 2065]]

GRad_BOOst

test_accuracy	0.868646
recall_test	0.904898
precision_test	0.617256
f1_test	0.733899
auc_test	0.882236
cm_test	[[7223, 1180], [200, 1903]]
train_accuracy	0.890849
precision_train	0.659499
recall_train	0.944418
f1_train	0.776652
auc_train	0.5
cm_train	[[22095, 3088], [352, 5981]]

Model Validation & Selection(SMOTE)

	Model Name	Train Accuracy	Test Accuracy	Precision Train	Precision Test	Recall Train	Recall test	ROC AUC Train	ROC AUC Test	AUC Train	AUC Test
0	LogisticRegression	0.84	0.83	0.83	0.54	0.85	0.81	0.84	0.82	0.838939	0.819453
1	DecisionTreeClassifier	1.00	0.88	1.00	0.68	1.00	0.72	1.00	0.82	1.000000	0.817953
2	RandomForestClassifier	1.00	0.92	1.00	0.77	1.00	0.83	1.00	0.89	1.000000	0.885250
3	GradientBoostingClassifier	0.94	0.91	0.93	0.76	0.95	0.84	0.94	0.89	0.940555	0.885723
4	XGBClassifier	0.94	0.91	0.93	0.76	0.95	0.84	0.94	0.89	0.938768	0.885544
5	KNeighborsClassifier	0.84	0.60	0.78	0.23	0.95	0.44	0.84	0.54	0.843744	0.539154
6	SVC	0.58	0.59	0.58	0.26	0.56	0.55	0.58	0.58	0.579597	0.575214

Model Validation & Selection(Classification Matrix)

Test Data


```
For LogisticRegression Confusion matrix is :  
[[6983 1420]  
 [ 404 1699]]  
For DecisionTreeClassifier Confusion matrix is :  
[[7701  702]  
 [ 590 1513]]  
For RandomForestClassifier Confusion matrix is :  
[[7893  510]  
 [ 355 1748]]  
For GradientBoostingClassifier Confusion matrix is :  
[[7849  554]  
 [ 342 1761]]  
For XGBClassifier Confusion matrix is :  
[[7842  561]  
 [ 341 1762]]  
For KNeighborsClassifier Confusion matrix is :  
[[5353 3050]  
 [1175  928]]  
For SVC Confusion matrix is :  
[[5040 3363]  
 [ 945 1158]]
```

Train Data

```
For LogisticRegression Confusion matrix is :  
[[20834  4349]  
 [ 3763 21420]]  
For DecisionTreeClassifier Confusion matrix is :  
[[25183    0]  
 [    0 25183]]  
For RandomForestClassifier Confusion matrix is :  
[[25183    0]  
 [    0 25183]]  
For GradientBoostingClassifier Confusion matrix is :  
[[23481  1702]  
 [ 1292 23891]]  
For XGBClassifier Confusion matrix is :  
[[23438  1745]  
 [ 1339 23844]]  
For KNeighborsClassifier Confusion matrix is :  
[[18531  6652]  
 [ 1218 23965]]  
For SVC Confusion matrix is :  
[[15080 10103]  
 [11071 14112]]
```


Model Validation & Selection (Hyperparameter tuned)

Accuracy Score 0.9237578526556254					
	precision	recall	f1-score	support	
0.0	0.95	0.95	0.95	8403	
1.0	0.81	0.81	0.81	2103	
accuracy			0.92	10506	
macro avg			0.88	10506	
weighted avg			0.92	10506	



**Random Forest
Classifier**

**XG Boost
Classifier**



Accuracy Score 0.73043974871502					
	precision	recall	f1-score	support	
0.0	0.99	0.67	0.80	8403	
1.0	0.43	0.99	0.59	2103	
accuracy			0.73	10506	
macro avg			0.71	10506	
weighted avg			0.88	10506	

Model Validation & Selection(continued)

Observation 1: As seen in the table above, DT classifier gives us a better precision value whereas Random Forest classifier is lagging behind in terms of precision, but from the business perspective we are looking for a better recall value, which is given to us by Random Forest classifier.


Observation 2: Gaussian NB has performed poorly in terms of precision and AUC value, whereas we see that DT Classifier, RT Classifier, GB Classifier & XGB Classifier have performed equally good in terms of F-1 Score and Test accuracy.



Model Validation & Selection(continued)

Observation 3: From the above observation we have come to the conclusion that we would choose our model from RF Classifier, GB Classifier or XGB Classifier.

Although, we see DT Classifier also giving great results but that may lead of overfitting of the data.



Model Validation & Selection(continued)


We had chosen XG Boost Classifier for our prediction and the best hyperparameters obtained are as below.

Best hyperparameters:

```
colsample_bylevel=1,  
colsample_bynode=1  
colsample_bytree=0.9,eta=0.3,gamma=0,  
learning_rate=0.1, max_delta_step=0, max_depth=5,  
min_child_weight=1, missing=None, n_estimators=100,  
n_jobs=1,nthread=4,num_boost_round=10  
objective='binary:logistic'  
random_state=0, reg_alpha=0,  
reg_lambda=1,scale_pos_weight=90
```



Model Validation & Selection(continued)



Challenges

- Huge chunk of data was to be handled keeping in mind not to miss anything which is even of little relevance.
- Feature selection was quite a challenge as our dataset had many futuristic features which had no relevance for initial detection of loan defaulter.

Futuristic Features

total_pymnt:- Payments received to date for total amount funded.

total_pymnt_inv:-Payments received on date for portion of total amount funded by investors.

total_rec_prncp:-Principal received to date.

total_rec_int:-Interest received to date.

total_rec_late_fee:-Late fees received to date.

recoveries:-post charge off gross recovery.

collection_recovery_fee:-post charge off collection fee.

last_pymnt_d:-Last month payment was received.

last_pymnt_amnt:-Last total payment amount received.

funded_amnt:-The total amount committed to that loan at that point in time.

funded_amnt_inv:-The total amount committed by investors for that loan at that time.

Q & A